



SPEED DATING
EXPERIMENT

What are the least desirable attributes in a male partner?

Speed Dating Experiment


- Columbia Business School.
- Experimental speed dating events from 2002-2004.
- 8378 dates.
- 4 minute dates with every other participant of the opposite sex.
- Participants were asked **if they would like to see their date again** (decision) and to **rate their date on 6 attributes** (0 to 10):
 1. Attractiveness
 2. Sincerity
 3. Intelligence
 4. Fun
 5. Ambition
 6. Shared Interests.



DATASET

- **Key Challenge: Dropping columns to make it more manageable.**
I decided to reduce it from 195 to less than 15 columns.
- **Missing values and holes in the data.**
If a column has many missing values,
I decided to drop it so it doesn't bias the result.
- **The column names were hard to read (amb7_2).**
I decided to rename some of them to something more readable to
make analyzing the data simpler.

How the speed dating rounds worked



1) There are two groups.

2) One group is women and the other is men.

3) The point of it all is to match every woman with every man for a short period of time so that **by the end, every one has gotten a chance to quickly know each other.**

The assumption here: it is possible to learn a lot about a person in a short period of time.

PARTICIPANTS

Total: 551

	MALE	FEMALE
Participants	277 (50,3%)	274 (49.7%)
Avg. dates	15.1	15.2
Match rate	20.2%	20.9%
Partner wanted a date	39.7%	48.6%

Number of unique participants

```
In [60]: f = data_df.loc[data_df.gender == 0]
m = data_df.loc[data_df.gender == 1]
print('The total number of dates is: {}'.format(len(data_df)))
print('The total number of unique participants is: {}'.format(len(data_df['iid'].unique())))
print('The number of female participants is: {}'.format(len(f['iid'].unique())))
print('The number of male participants is: {}'.format(len(m['iid'].unique())))
```

```
The total number of dates is: 8378
The total number of unique participants is: 551
The number of female participants is: 274
The number of male participants is: 277
```

Number of match per gender

```
In [65]: data_m = data_df[data_df['gender']==1].groupby('iid').sum()
match_m = data_m['match']
g = plt.figure(figsize=(5,5))
g = plt.hist(match_m, range(15))
g = plt.xlabel('Male Matches', fontsize=14)

# Number of dates males
dates_male = data_df[data_df.gender == 1].groupby('iid').apply(len)

# The of matches males
matches_m = data_df[data_df.match == 1]
matches_male = matches_m[matches_m.gender == 1].groupby('iid').apply(len)


#Male match percentage
mmp = (matches_male / dates_male).mean() * 100.0
mmp
print('The avg. dates per male is: %s' %(dates_male.mean()))
print('The match percentage for males is : %s' % (mmp))

# Date? decision of partner == 1
partner_yes_M = data_df[data_df.dec_o == 1]
partner_syes_M = partner_yes_M[partner_yes_M.gender == 1].groupby('iid').apply(len)
pyp_M = ((partner_syes_M/dates_male).mean())*100
print('Female partner said yes %s percent of the times' % (pyp_M))
```

The avg. dates per male is: 15.1407942238

The match percentage for males is : 20.2888324872

Female partner said yes 39.7366771615 percent of the times



```

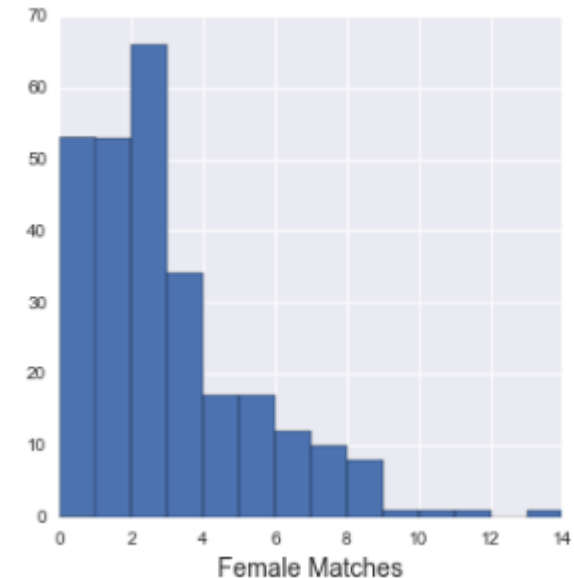
In [66]: data_f = data_df[data_df['gender']==0].groupby('iid').sum()
match_f = data_f['match']
g = plt.figure(figsize=(5,5))
g = plt.hist(match_f, range(15))
g = plt.xlabel('Female Matches', fontsize=14)

# Number of dates Females
dates_female = data_df[data_df.gender == 0].groupby('iid').apply(len)
# The of matches Females
matches = data_df[data_df.match == 1]
matches_female = matches[matches.gender == 0].groupby('iid').apply(len)
#Female match percentage
fmp = (matches_female / dates_female).mean() * 100.0
fmp
print('The avg. dates per female is: %s' %(dates_female.mean()))
print('The match percentage for females is : %s' % (fmp))

# Date? decision of partner == 1
partner_yes_F = data_df[data_df.dec_o == 1]
partner_syes_F = partner_yes_F[partner_yes_F.gender == 0].groupby('iid').apply(len)
pyp_F = ((partner_syes_F/dates_female).mean())*100
print('Male partner said yes %s percent of the times' % (pyp_F))

```


The avg. dates per female is: 15.2700729927
 The match percentage for females is : 20.9103753144
 Male partner said yes 48.6359814379 percent of the times



Participant's age distribution


```
data_unique = data_df.groupby('iid').mean()  
data_unique.age.plot(kind='hist', bins=6)  
print("participant's age distribution")
```

participant's age distribution




```
data_unique_f = data_df[data_df['gender']==0].groupby('iid').mean()  
data_unique_f.age.plot(kind='hist', bins=6)  
print("Females age distribution")
```

Females age distribution




```
data_unique_m = data_df[data_df['gender']==1].groupby('iid').mean()  
data_unique_m.age.plot(kind='hist', bins=6)  
print("Males age distribution")
```

Males age distribution



Participant's activities (Interest)

```
activities_interested=['sports','tvsports', 'dining', 'museums', 'art', 'hiking', 'reading', 'tv', 'shopping',  
                      'theater','movies', 'concerts', 'clubbing', 'gaming', 'yoga']  
temp = data_df.groupby(['gender']).mean()[activities_interested].values  
  
g = plt.figure(figsize=(8,8))  
g = plt.barh(np.arange(0,2*temp.shape[1],2)+0.2,temp[1,:], height=0.5,color=[0,0,1],alpha=0.5,label='Male')  
g = plt.barh(np.arange(0,2*temp.shape[1],2)-0.2,temp[0,:], height=0.5,color=[1,0,1],alpha=0.5,label='Female')  
g = plt.yticks(np.arange(0,2*temp.shape[1],2)+0.2,activities_interested,fontsize=14)  
g = plt.ylim(-1,2*temp.shape[1]+1)  
g = plt.legend(loc=0,fontsize=14)
```



Getting Dummies

```
In [67]: dummy_gender = pd.get_dummies(data_df['gender'], prefix='gender')
print dummy_gender.head()
```

```
gender_0  gender_1
0         1.0      0.0
1         1.0      0.0
2         1.0      0.0
3         1.0      0.0
4         1.0      0.0
```

Columns to keep

```
cols_to_keep = ['dec_o', 'age', 'attr_o', 'sinc_o', 'intel_o', 'fun_o', 'amb_o', 'shar_o']
data_b = data_df[cols_to_keep].join(dummy_gender.ix[:, 'gender':])
data_b.describe()
```

	dec_o	age	attr_o	sinc_o	intel_o	fun_o	amb_o	shar_o	gender_0	gender_1
count	8378.000000	8283.000000	8166.000000	8091.000000	8072.000000	8018.000000	7656.000000	7302.000000	8378.000000	8378.000000
mean	0.419551	26.358928	6.190411	7.175256	7.369301	6.400599	6.778409	5.474870	0.499403	0.500597
std	0.493515	3.566763	1.950305	1.740575	1.550501	1.954078	1.794080	2.156163	0.500029	0.500029
min	0.000000	18.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	24.000000	5.000000	6.000000	6.000000	5.000000	6.000000	4.000000	0.000000	0.000000
50%	0.000000	26.000000	6.000000	7.000000	7.000000	7.000000	7.000000	6.000000	0.000000	1.000000
75%	1.000000	28.000000	8.000000	8.000000	8.000000	8.000000	8.000000	7.000000	1.000000	1.000000
max	1.000000	55.000000	10.500000	10.000000	10.000000	11.000000	10.000000	10.000000	1.000000	1.000000

```
# intercept
data_b['intercept'] = 1.0
data_b.head()
```

	dec_o	age	attr_o	sinc_o	intel_o	fun_o	amb_o	shar_o	gender_0	gender_1	intercept
0	0	21.0	6.0	8.0	8.0	8.0	8.0	6.0	1.0	0.0	1.0
1	0	21.0	7.0	8.0	10.0	7.0	7.0	5.0	1.0	0.0	1.0
2	1	21.0	10.0	10.0	10.0	10.0	10.0	10.0	1.0	0.0	1.0
3	1	21.0	7.0	8.0	9.0	8.0	9.0	8.0	1.0	0.0	1.0
4	1	21.0	8.0	7.0	9.0	6.0	9.0	7.0	1.0	0.0	1.0

Dropping data points with missing data

```
data_b.rename(columns={'dec_o': 'decision', 'attr_o': 'attractive', 'sinc_o': 'sincere', 'intel_o': 'intelligent', 'fun_o': 'fun',
                      'amb_o': 'ambitious', 'shar_o': 'shared interests'}, inplace=True)
```

```
data_c = data_b.dropna()
data_c.isnull().sum()
```

```
decision      0
age           0
attractive    0
sincere       0
intelligent   0
fun           0
ambitious     0
shared interests 0
gender_0      0
gender_1      0
intercept     0
dtype: int64
```

```
len(data_c)
```

```
6959
```

```
data_c.isnull().count()
```

```
decision      6959  
age           6959  
attractive    6959  
sincere       6959  
intelligent   6959  
fun           6959  
ambitious     6959  
shared interests 6959  
gender_0      6959  
gender_1      6959  
intercept     6959  
dtype: int64
```

General Dataset descriptive measure

```
data_c.describe()
```

	decision	age	attractive	sincere	intelligent	fun	ambitious	shared interests	gender_0	gender_1	inte
count	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000
mean	0.429947	26.318437	6.183561	7.162954	7.361690	6.395315	6.759089	5.460052	0.506538	0.493462	1.0
std	0.495104	3.564386	1.949638	1.745162	1.559914	1.959143	1.797901	2.149901	0.499993	0.499993	0.0
min	0.000000	18.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.0
25%	0.000000	24.000000	5.000000	6.000000	6.000000	5.000000	6.000000	4.000000	0.000000	0.000000	1.0
50%	0.000000	26.000000	6.000000	7.000000	7.000000	7.000000	7.000000	6.000000	1.000000	0.000000	1.0
75%	1.000000	28.000000	8.000000	8.000000	8.000000	8.000000	8.000000	7.000000	1.000000	1.000000	1.0
max	1.000000	55.000000	10.000000	10.000000	10.000000	11.000000	10.000000	10.000000	1.000000	1.000000	1.0

```
data_c = data_c.replace(11, 10)  
data_c.describe()
```

	decision	age	attractive	sincere	intelligent	fun	ambitious	shared interests	gender_0	gender_1	inte
count	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000	6959.000000
mean	0.429947	26.318437	6.183561	7.162954	7.361690	6.395172	6.759089	5.460052	0.506538	0.493462	1.0
std	0.495104	3.564386	1.949638	1.745162	1.559914	1.958842	1.797901	2.149901	0.499993	0.499993	0.0
min	0.000000	18.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.0
25%	0.000000	24.000000	5.000000	6.000000	6.000000	5.000000	6.000000	4.000000	0.000000	0.000000	1.0
50%	0.000000	26.000000	6.000000	7.000000	7.000000	7.000000	7.000000	6.000000	1.000000	0.000000	1.0
75%	1.000000	28.000000	8.000000	8.000000	8.000000	8.000000	8.000000	7.000000	1.000000	1.000000	1.0
max	1.000000	55.000000	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000	1.000000	1.000000	1.0

CORRELATION

DATASET

```
data_c.corr()
```

	decision	age	attractive	sincere	intelligent	fun	ambitious	shared interests	gender_0	gender_1	intercept
decision	1.000000	-0.046645	0.487717	0.207160	0.214155	0.411843	0.184109	0.400070	0.117529	-0.117529	NaN
age	-0.046645	1.000000	-0.047709	0.004041	0.033008	-0.035223	0.019968	0.005233	-0.071570	0.071570	NaN
attractive	0.487717	-0.047709	1.000000	0.406055	0.388974	0.590472	0.359268	0.490608	0.129130	-0.129130	NaN
sincere	0.207160	0.004041	0.406055	1.000000	0.667933	0.507764	0.464358	0.398944	0.041191	-0.041191	NaN
intelligent	0.214155	0.033008	0.388974	0.667933	1.000000	0.500992	0.629279	0.401784	-0.057852	0.057852	NaN
fun	0.411843	-0.035223	0.590472	0.507764	0.500992	1.000000	0.493640	0.617335	0.058479	-0.058479	NaN
ambitious	0.184109	0.019968	0.359268	0.464358	0.629279	0.493640	1.000000	0.434890	-0.098770	0.098770	NaN
shared interests	0.400070	0.005233	0.490608	0.398944	0.401784	0.617335	0.434890	1.000000	0.029991	-0.029991	NaN
gender_0	0.117529	-0.071570	0.129130	0.041191	-0.057852	0.058479	-0.098770	0.029991	1.000000	-1.000000	NaN
gender_1	-0.117529	0.071570	-0.129130	-0.041191	0.057852	-0.058479	0.098770	-0.029991	-1.000000	1.000000	NaN
intercept	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

MULTICOLLINEARITY ??

Dataset (Female only)

```
subF = data_c[(data_c['gender_0']== 1)]
del subF['gender_1']
subF.head()
```

	decision	age	attractive	sincere	intelligent	fun	ambitious	shared interests	gender_0	intercept
0	0	21.0	6.0	8.0	8.0	8.0	8.0	6.0	1.0	1.0
1	0	21.0	7.0	8.0	10.0	7.0	7.0	5.0	1.0	1.0
2	1	21.0	10.0	10.0	10.0	10.0	10.0	10.0	1.0	1.0
3	1	21.0	7.0	8.0	9.0	8.0	9.0	8.0	1.0	1.0
4	1	21.0	8.0	7.0	9.0	6.0	9.0	7.0	1.0	1.0

Analysis

```
data_fem.corr()
```

[illegible]

```
# covariates
train_cols = data_fem.columns[1:]
print train_cols
```

```
Index([u'age', u'attractive', u'sincere', u'intelligent', u'fun', u'ambitious',
      u'shared interests', u'intercept'],
      dtype='object')
```

```
# Fit the model
```

```
logit_fem = sm.Logit(data_fem['decision'], data_fem[train_cols])
result_F = logit_fem.fit()
```

```
Optimization terminated successfully.
Current function value: 0.501235
Iterations 6
```

```
print result_F.summary()
```

Logit Regression Results

```
=====
Dep. Variable:          decision    No. Observations:          3525
Model:                  Logit      Df Residuals:            3517
Method:                 MLE        Df Model:                  7
Date:                  Mon, 08 Aug 2016    Pseudo R-squ.:          0.2765
Time:                  11:18:26    Log-Likelihood:         -1766.9
converged:              True        LL-Null:                -2442.2
                                   LLR p-value:                1.764e-287
=====
```

	coef	std err	z	P> z	[95.0% Conf. Int.]
age	-0.0047	0.011	-0.412	0.681	-0.027 0.018
attractive	0.6912	0.033	20.675	0.000	0.626 0.757
sincere	-0.1585	0.037	-4.274	0.000	-0.231 -0.086
intelligent	-0.0356	0.043	-0.826	0.409	-0.120 0.049
fun	0.2593	0.035	7.502	0.000	0.192 0.327
ambitious	-0.1574	0.034	-4.604	0.000	-0.224 -0.090
shared interests	0.2671	0.026	10.141	0.000	0.215 0.319
intercept	-5.1709	0.398	-12.981	0.000	-5.952 -4.390

```
=====
```

MULTICOLLINEARITY!!!

Dataset (Male only)

```
subM = data_c[(data_c['gender_1']== 1)]
del subM['gender_0']
subM.tail()
```

	decision	age	attractive	sincere	intelligent	fun	ambitious	shared interests	gender_1	intercept
8373	1	25.0	10.0	5.0	3.0	2.0	6.0	5.0	1.0	1.0
8374	0	25.0	6.0	3.0	7.0	3.0	7.0	2.0	1.0	1.0
8375	0	25.0	2.0	1.0	2.0	2.0	2.0	1.0	1.0	1.0
8376	1	25.0	5.0	7.0	5.0	5.0	3.0	6.0	1.0	1.0
8377	1	25.0	8.0	8.0	7.0	7.0	7.0	7.0	1.0	1.0

```
data_male.corr()
```

[illegible]

```
# covariates
train_cols = data_male.columns[1:]
print train_cols
```

```
Index([u'age', u'attractive', u'sincere', u'intelligent', u'fun', u'ambitious',
      u'shared interests', u'intercept'],
      dtype='object')
```

```
# Fit the model
```

```
logit_male = sm.Logit(data_male['decision'], data_male[train_cols])
result_M = logit_male.fit()
```

```
Optimization terminated successfully.
      Current function value: 0.501609
      Iterations 6
```

```
print result_M.summary()
```

Logit Regression Results

```
=====
Dep. Variable:          decision    No. Observations:          3434
Model:                  Logit      Df Residuals:              3426
Method:                  MLE       Df Model:                  7
Date:                   Mon, 08 Aug 2016    Pseudo R-squ.:          0.2394
Time:                   17:15:37          Log-Likelihood:         -1722.5
converged:               True          LL-Null:                 -2264.7
                                   LLR p-value:              7.367e-230
=====
```

	coef	std err	z	P> z	[95.0% Conf. Int.]
age	-0.0223	0.013	-1.784	0.074	-0.047 0.002
attractive	0.4045	0.029	13.748	0.000	0.347 0.462
sincere	-0.0943	0.035	-2.686	0.007	-0.163 -0.025
intelligent	0.1300	0.045	2.894	0.004	0.042 0.218
fun	0.2697	0.034	8.026	0.000	0.204 0.336
ambitious	-0.1654	0.034	-4.795	0.000	-0.233 -0.098
shared interests	0.2794	0.027	10.415	0.000	0.227 0.332
intercept	-4.9712	0.425	-11.699	0.000	-5.804 -4.138

```
=====
```

MULTICOLLINEARITY!!!

NEXT STEP

- **Determine new** “Categories of features”
- **Create a new** variables with the mean of the variables in that group
- or just keep one of the variables.

